

Introducing **Symbolic** Execution

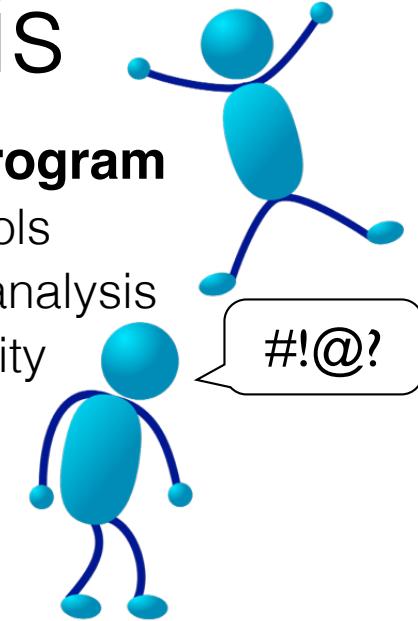


Software has **bugs**

- To **find them**, we use **testing** and **code reviews**
- But some **bugs are still missed**
 - *Rare features*
 - *Rare circumstances*
 - *Nondeterminism*

Static analysis

- Can **analyze all possible runs of a program**
 - An explosion of interesting ideas and tools
 - Commercial companies sell, use static analysis
 - Great potential to improve software quality
- But: **Can it find deep, difficult bugs?**
 - Our experience: yes, but **not often**
 - Commercial viability implies you must deal with developer confusion, false positives, error management,..
 - This means that companies specifically aim to keep the **false positive rate down**
 - They often do this **by purposely missing bugs**, to keep the analysis simpler



One issue: Abstraction

- Abstraction lets us model all possible runs
 - But **abstraction introduces conservatism**
- ***-sensitivities add precision**, to deal with this
 - * = **flow-, context-, path-, etc.**
 - But **more precise abstractions are more expensive**
 - Challenges scalability
 - Still have false alarms or missed bugs
- **Static analysis abstraction ≠ developer abstraction**
 - Because the developer didn't have them in mind

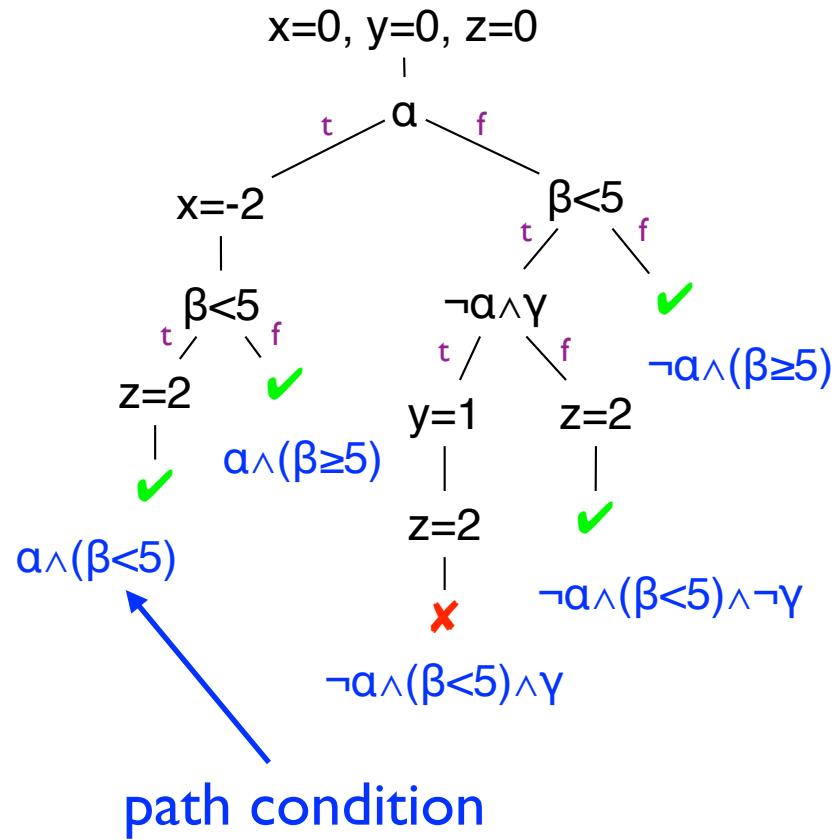
Symbolic execution

A middle ground

- Testing works: reported bugs are real bugs
 - But, **each test** only explores **one possible execution**
 - `assert(f(3) == 5)`
 - In short, **complete**, but **not sound**
 - We *hope* test cases generalize, but no guarantees
- **Symbolic execution generalizes testing**
 - “More sound” than testing
 - Allows unknown symbolic variables `a` in evaluation
 - `y = a; assert(f(y) == 2*y-1);`
 - If execution path depends on unknown, conceptually fork symbolic executor
 - `int f(int x) { if (x > 0) then return 2*x - 1; else return 10; }`

Symbolic execution example

```
1. int a = α, b = β, c = γ;  
2.                                // symbolic  
3. int x = 0, y = 0, z = 0;  
4. if (a) {  
5.     x = -2;  
6. }  
7. if (b < 5) {  
8.     if (!a && c) { y = 1; }  
9.     z = 2;  
10.}  
11.assert(x+y+z != 3)
```



Insight

- Each **symbolic execution path** stands for **many** actual **program runs**
 - In fact, exactly the set of runs whose concrete values satisfy the path condition
- Thus, we can **cover a lot more of the program's execution space than testing**
- Viewed **as a static analysis, symbolic execution** is
 - **Complete**, but not sound (usually doesn't terminate)
 - **Path, flow, and context sensitive**